

Introduction to Machine Learning with Python

David Schaupp | WS2025

Content

Introduction to Machine Learning:

- What is Machine Learning?
- Machine Learning Project Checklist

Intro's:

- Colab
- Python
- Numpy
- Matplotlib

Supervised Learning:

- Classification (Binary|Multiclass)
- **Regression**
- Support Vector Machines
- Decision Trees (Random Forest)

Unsupervised Learning:

- Dimensionality Reduction
- Clustering (k-means, DBSCAN)

Performances Measures:

- Accuracy
- Precision
- Recall
- F1-Score
- ROC Curve
- Confusion Matrix



Regression vs Classification

Comparison:

- Type of Prediction

Regression: Predicts continuous values

Classification: Predicts discrete values

- Example

Regression: Predicting the price of a house in Vienna

Classification: Predicting if a house is in Vienna or not

- Task of Algorithm

Regression: Mapping input value (x) with continuous output variable (y)

Classification: Mapping the input value of x with the discrete output variable of y

- Problem Solution

Regression: Solves regression problems such as house price predictions and weather predictions

Classification: Solves classification problems like identifying spam e-mails, spotting cancer cells and speech recognition

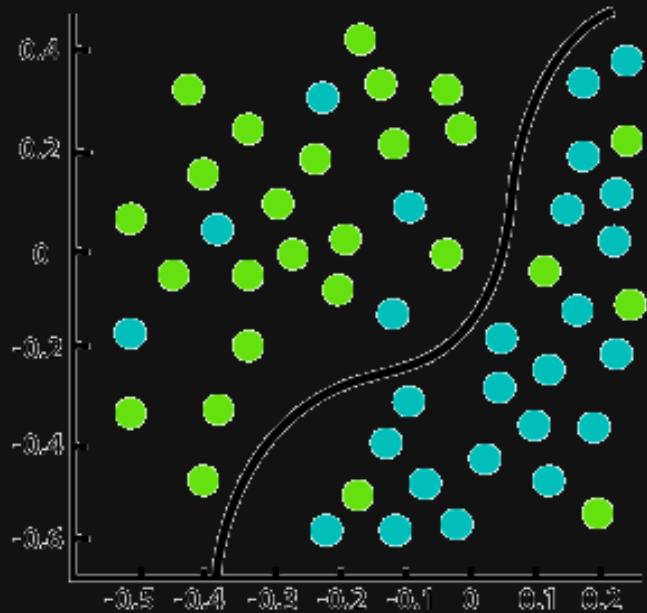
- Further Division

Regression: Can be divided into Linear and Non-linear Regression

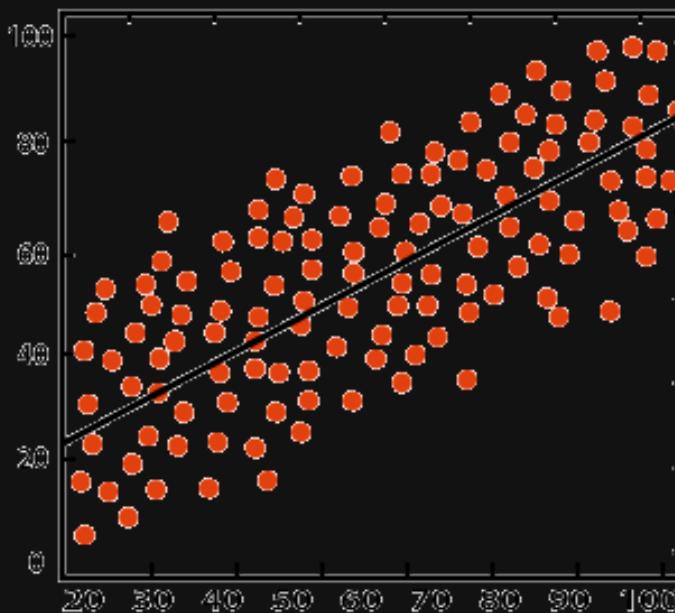
Classification: Can be divided into Binary Classifiers and Multi-class Classifiers



Regression vs Classification



Classification



Regression



Linear Regression

Problem:

- Predict the price of a house, given the size of the house

Real Estate Agent Question:

- How much should I sell a 1250 sqft house for?

Solution:

- Fit a "straight line" (linear model) through the training data

Result:

- Intersect the line with 1250 sqft on x-axis
- Get estimated price on y-axis



Linear Regression - Terminology

Notation:

- Dataset: Data used to train the model
- x : "input" variable/features (size in feet²)
- y : "output" variable/target variable
(price in \$1000's)
- m : Number of training examples (size of the dataset)
- (x, y) : One training example
 - $(x^{(i)}, y^{(i)})$: i -th training example
 - First training example: $(x^{(1)}, y^{(1)}) = (2104, 400)$

size in feet ²	price in \$1000's
2104	400
1416	232
1534	315
852	178
...	...
3210	870



Linear Regression - Training Workflow

Training set:

- features x
- targets y

Learning algorithm:

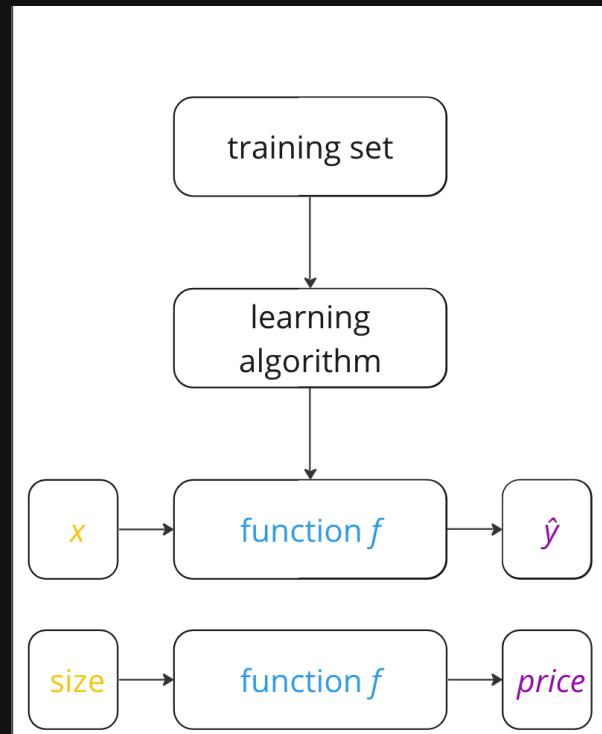
- Plug in x and y to learn patterns

Algorithm produces function f :

- Takes new input x
- Outputs prediction/estimate " \hat{y} "

House price example:

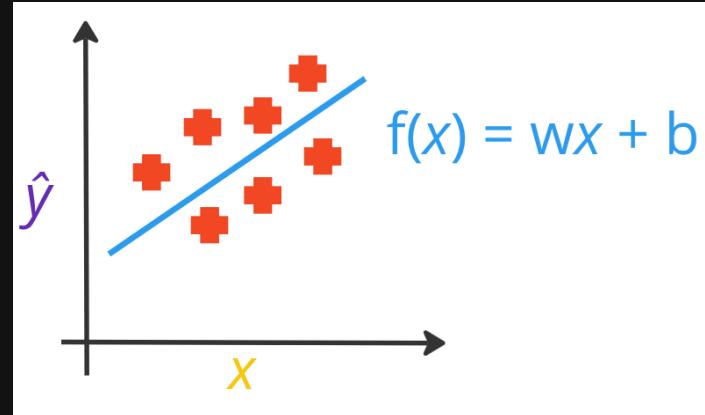
- Plug size (x) into function/model
- Get price (\hat{y}) from the function/model



Linear Regression - How to represent f?

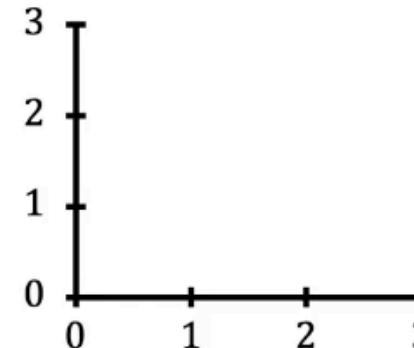
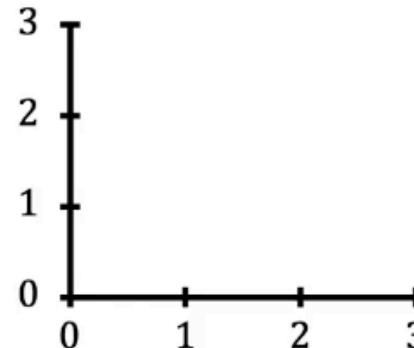
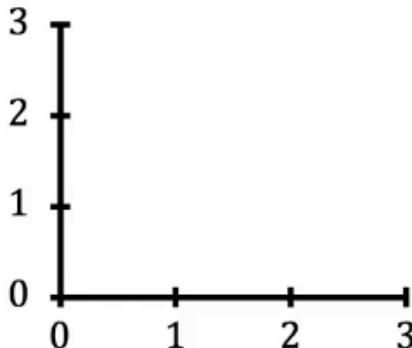
Model f :

- Linear function
- $f(x) = wx + b$
- w : weight
- b : bias
- w and b are parameters of the model

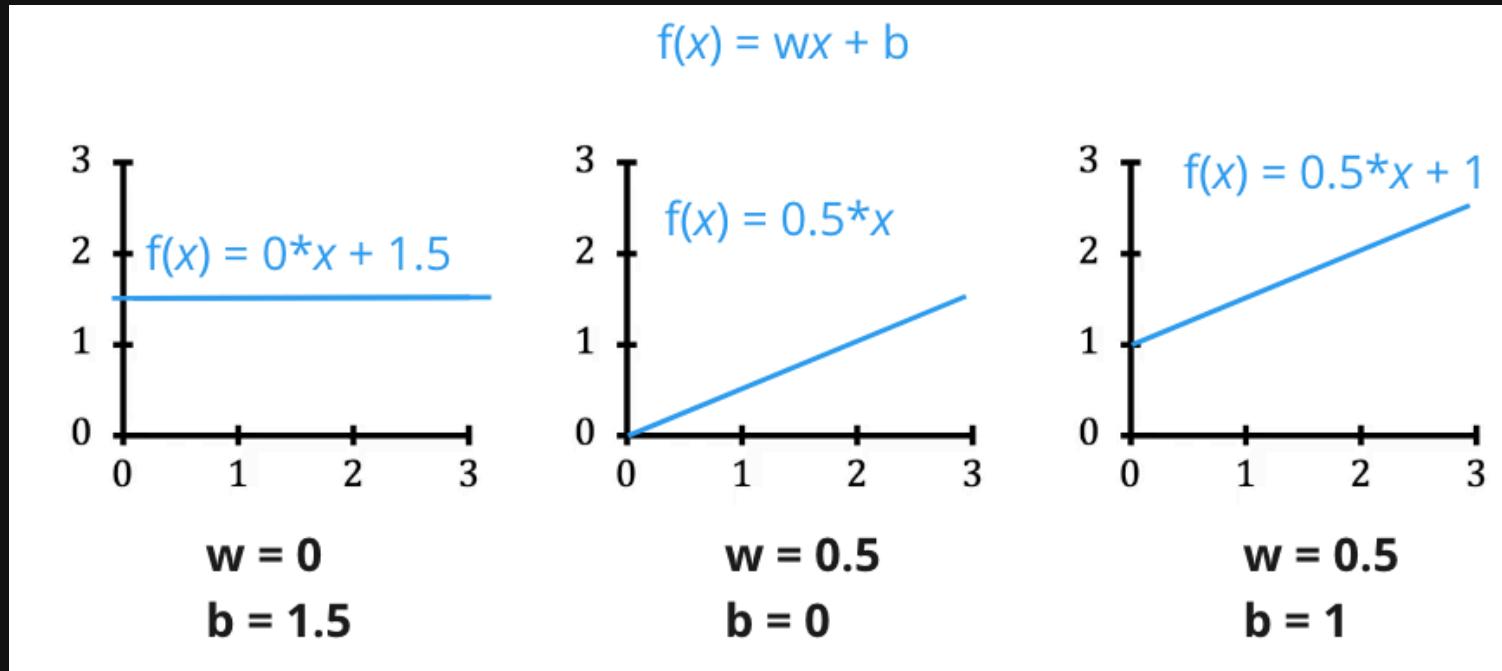


Linear Regression - How to represent f: Example

$$f(x) = wx + b$$



Linear Regression - How to represent f: Example



Linear Regression - How to train?

Model training:

- Setting the parameters w and b so that the model best fits the training data

Measure the fit of the model:

- Cost function

Cost function for linear regression:

- Squared Error Cost Function (Loss Function)

$$J_{(w,b)} = \frac{1}{m} \sum_m^{i=1} (f_{w,b}(x^{(i)}) - y^{(i)})^2$$

Linear Regression - Squared Error (Construction)

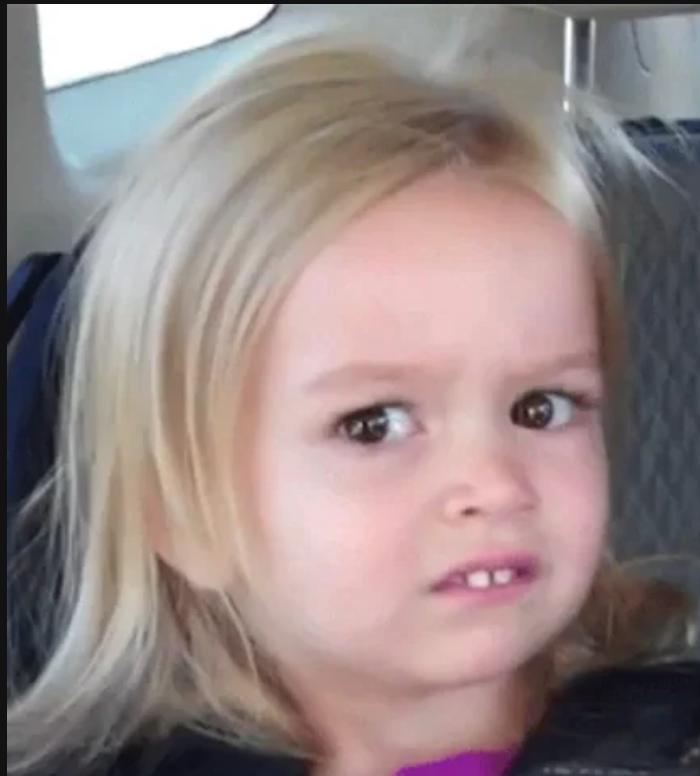
$$J_{(w,b)} = \frac{1}{m} \sum_m^{i=1} (f_{w,b}(x^{(i)}) - y^{(i)})^2$$

Formula Construction Steps:

1. Measure error: difference between prediction $y\hat{(i)}$ and actual $y^{(i)}$
2. Square the error: make it positive
3. Sum across all training examples: total error
4. Divide by m : average error
5. Divide by 2: mathematical convenience for gradient calculation



Enough maths (for now?)



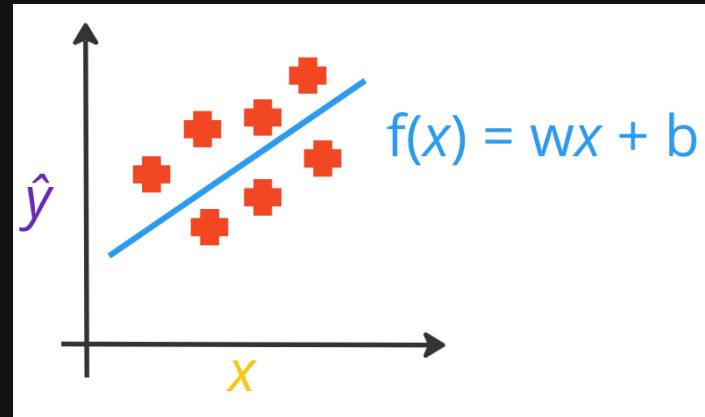
Linear Regression - Training

How to estimate the best parameters w and b for an given dataset?

- Minimize the cost function

How to minimize the cost function?

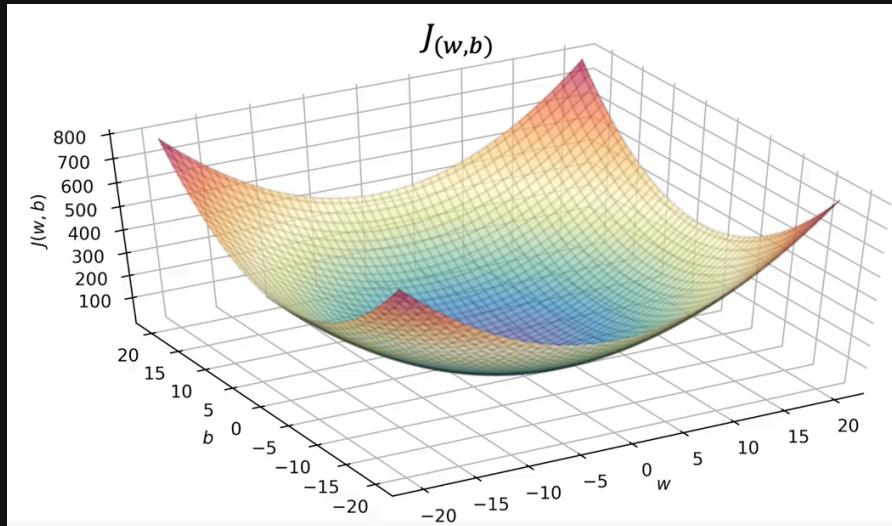
- Gradient descent



Linear Regression - Training

Squared Error Cost function:

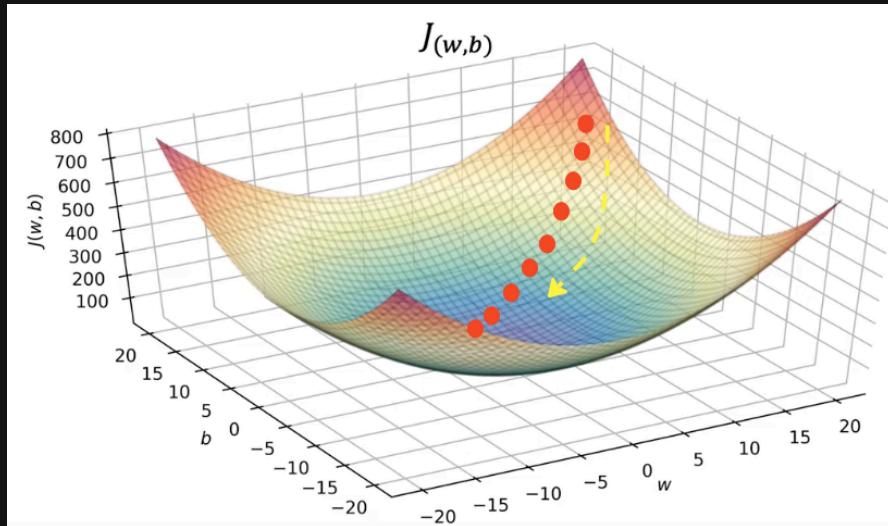
- 3D surface plot: axes labeled w and b , height is cost J
- Convex function (bowl shaped) - only one global minimum
- Gradient descent will find the global minimum
- Goal: Find parameters w and b that minimize $J(w,b)$



Linear Regression - Training

Gradient Descent:

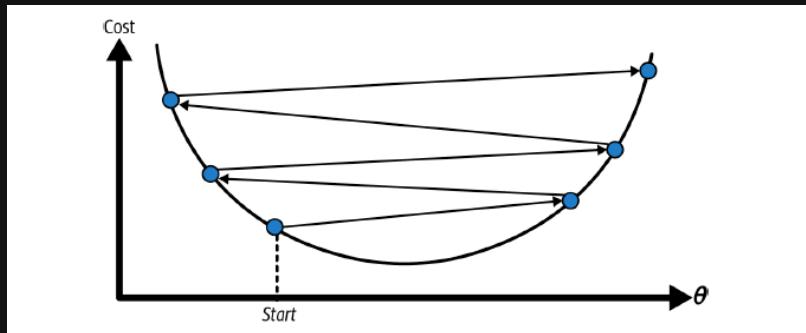
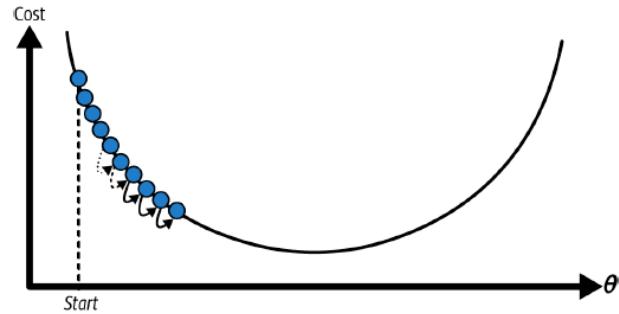
- Optimization algorithm to find optimal solution
- General idea: tweak parameters iteratively to minimize the cost function
- Steps:
 - Start with random parameters of w and b (random initialization)
 - Predict $y\text{-hat}$ for each training example and calculate the error
 - Calculate the gradient of the cost function
 - Update the parameters in the opposite direction of the gradient
 - Repeat until convergence (until the gradient is zero)



Linear Regression - Training

Learning Rate (Step Size):

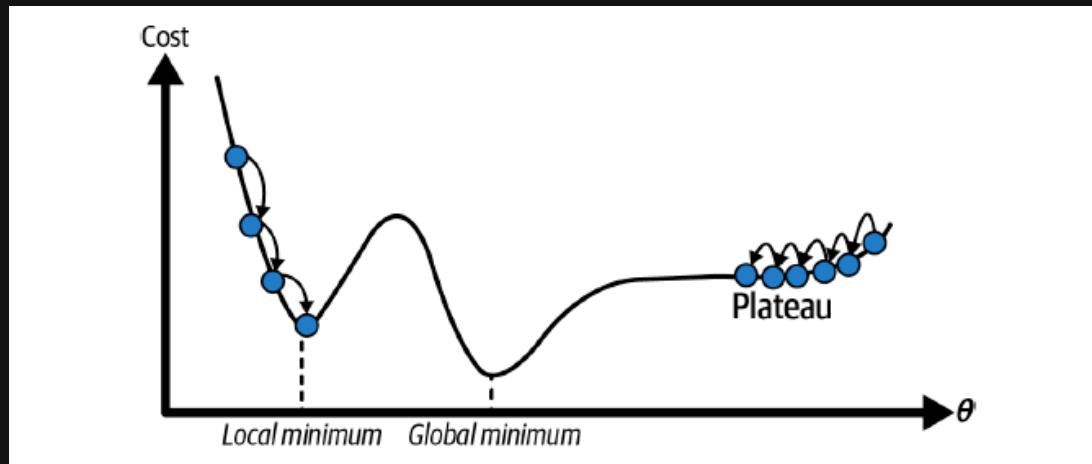
- Critical hyperparameter
- Too small:
 - Slow convergence, long training
- Too large:
 - Risk of overshooting minimum



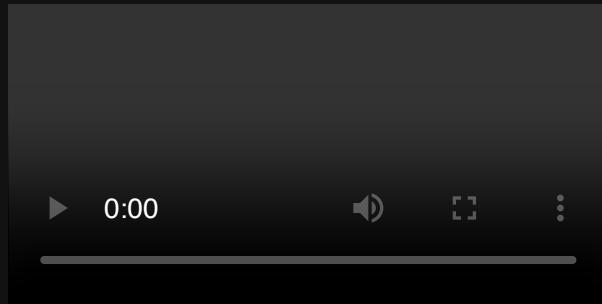
Linear Regression - Training

Gradient Descent:

- Not all functions look like nice bowls
- There are ridges, plateaus, etc.
- Convergence is not that easy
- If the random initialization starts the algorithm on the left:
 - The algorithm will converge to the local minimum
- If the random initialization starts the algorithm on the right:
 - It will take long time to cross the plateau
 - If training stops too early, the global minimum is not found



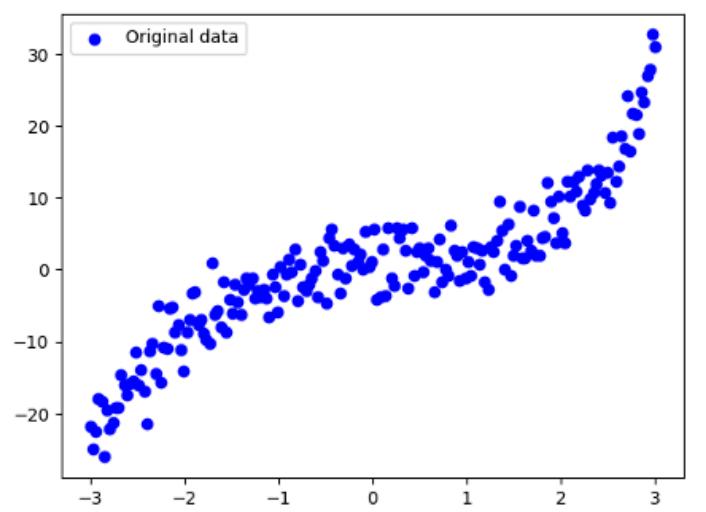
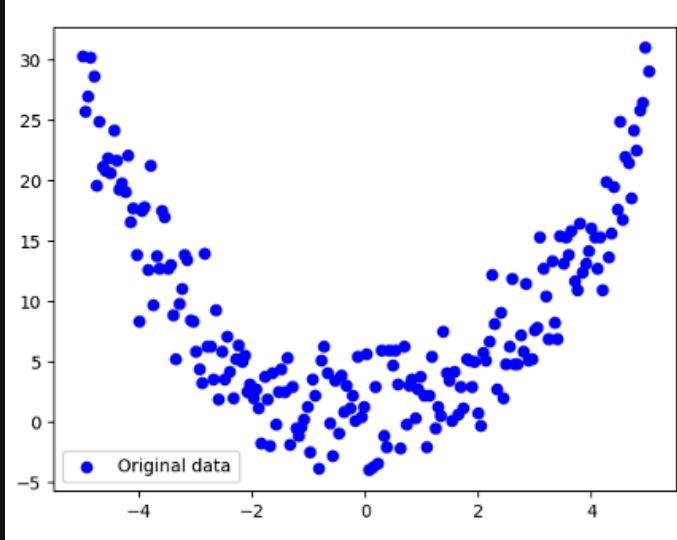
Linear Regression - Training



19

Polynomial Regression

How to Deal with such data:



Polynomial Regression

Solution: Add polynomial degrees to linear regression

- Remember: $y = wx + b$ (Linear Model - Degree 1)

For non-linear data:

- Degree 2 (Quadratic): $y = w_1x + w_2x^2 + b$

Examples:

- Degree 3 (Cubic): $y = w_1x + w_2x^2 + w_3x^3 + b$
- And higher degrees...

Key Challenge:

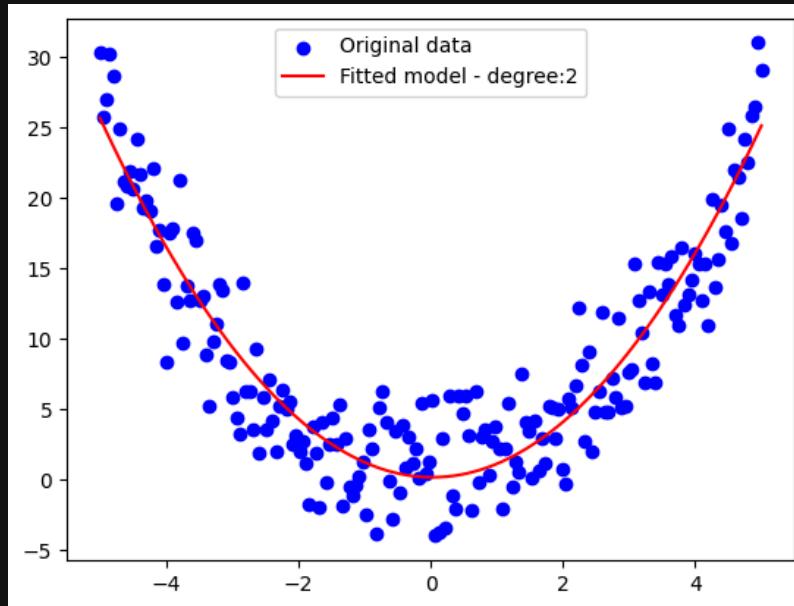
- Right degree = good fit
- Too high degree = overfitting



Polynomial Regression

Quadratic Model with Degree: 2

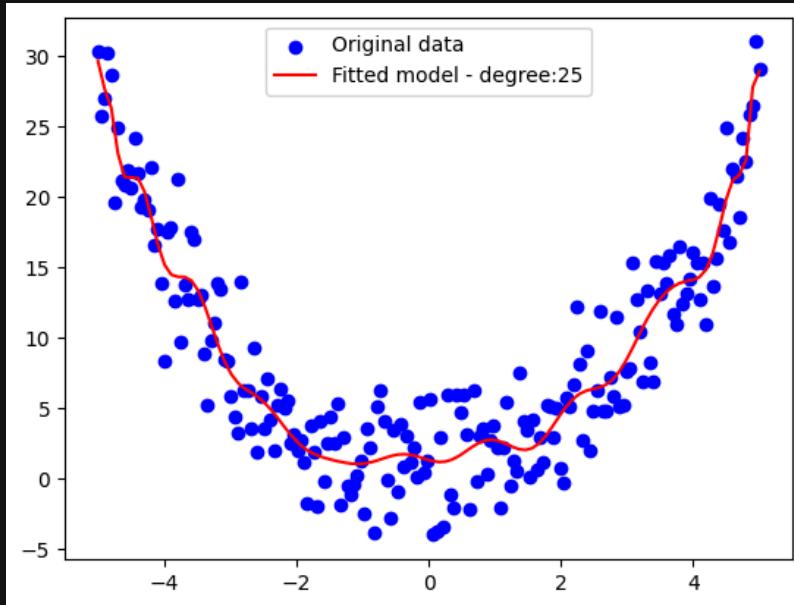
- Perfect fit for the data



Polynomial Regression

Quadratic Model with Degree: 25

- Model is Overfitting



Evaluation Metrics for Regression

- Mean Absolute Error (MAE):
 - Measures average prediction error
 - When to use: Best for understanding typical prediction error
 - Example Dataset: Balanced datasets without extreme outliers
- Mean Squared Error (MSE):
 - Emphasizes larger errors
 - When to use: Ideal when large errors need to be heavily penalized
 - Example Dataset: Financial predictions where large deviations are costly
- Root Mean Squared Error (RMSE):
 - Similar to MSE, but better for data with outliers
 - When to use: Use when outliers are significant in the dataset
 - Example Dataset: Environmental data with occasional extreme values
- **R² Score:**
 - Indicates fit quality (range from 0 to 1)
 - When to use: For overall fit assessment and model comparison
 - Example Dataset: General-purpose

$$MAE = \frac{1}{N} \sum_{i=1}^N |y_i - \hat{y}|$$

$$MSE = \frac{1}{N} \sum_{i=1}^N (y_i - \hat{y})^2$$

$$RMSE = \sqrt{MSE} = \sqrt{\frac{1}{N} \sum_{i=1}^N (y_i - \hat{y})^2}$$

$$R^2 = 1 - \frac{\sum (y_i - \hat{y})^2}{\sum (y_i - \bar{y})^2}$$

Where,

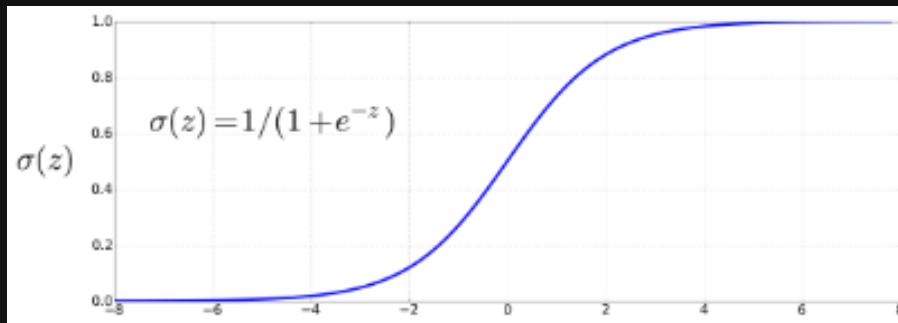
\hat{y} – predicted value of y

\bar{y} – mean value of y



Logistic Regression

- Classification Algorithm:
 - Binary Classification (0 or 1)
- Take Linear Regression Formula:
 - $y = wx + b$
- Plug it into the Sigmoid Function:
 - sigmoid = $1 / (1 + e^{-y})$
- Outputs are between 0 and 1:
 - 0.5 is the threshold
 - If sigmoid > 0.5 → 1
 - If sigmoid < 0.5 → 0



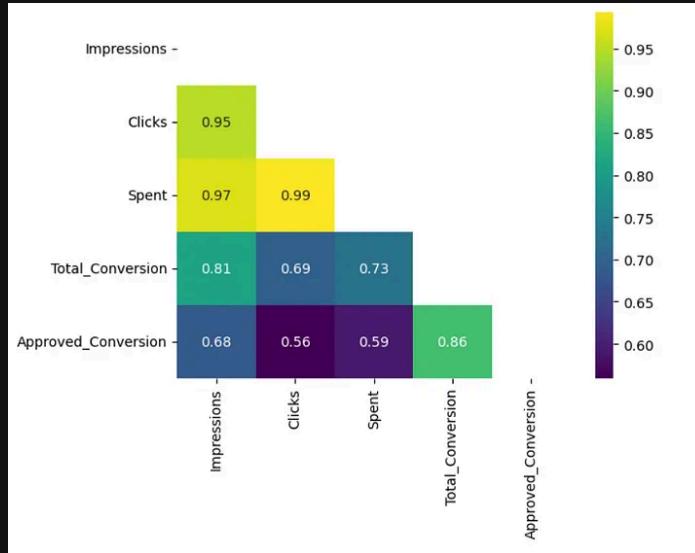
Feature Selection

- Feature Selection: Identifying relevant features to improve model performance
- Example: Predicting if a car should be crushed
 - Relevant: Model, Year, Miles driven
 - Irrelevant: Owner's name
- Methods:
 - Use statistical measures (e.g., Pearson Correlation)
- Approaches:
 - Filter: Statistical methods
 - Wrapper: Test different subsets
 - Embedded: Integrate into model training



Pearson Correlation Coefficient

- Pearson Correlation:
 - Measures linear relationship strength
 - Range: -1 (negative) to +1 (positive), 0 = no correlation
- Marketing Metrics Example:
 - Spent & Clicks: 0.99 (very strong)
 - Clicks & Impressions: 0.95 (strong)
 - Approved Conversion & Total Conversion: 0.86 (strong)
- Application:
 - Focus budget on metrics with strong correlations
 - Identify which features truly impact your target



Feature Scaling

- Example:
 - Wine chemical analysis dataset
 - Class based on alcohol and malic acid content
- Features are in different scales:
 - Alcohol: 11.03 - 14.83
 - Malic Acid: 0.74 - 5.80
- Features with larger scales will dominate the learning process
- By scaling features to the same range, we avoid this problem

Class	Alcohol	Malic
1	13.20	1.78
1	13.16	2.36
1	14.37	1.95
1	13.24	2.59
1	14.20	1.76



Feature Scaling - Methods

Normalization (Min-Max):

- Scales features to range [0, 1]
- Formula: $x' = (x - \min) / (\max - \min)$

$$x' = \frac{x - \min(x)}{\max(x) - \min(x)}$$

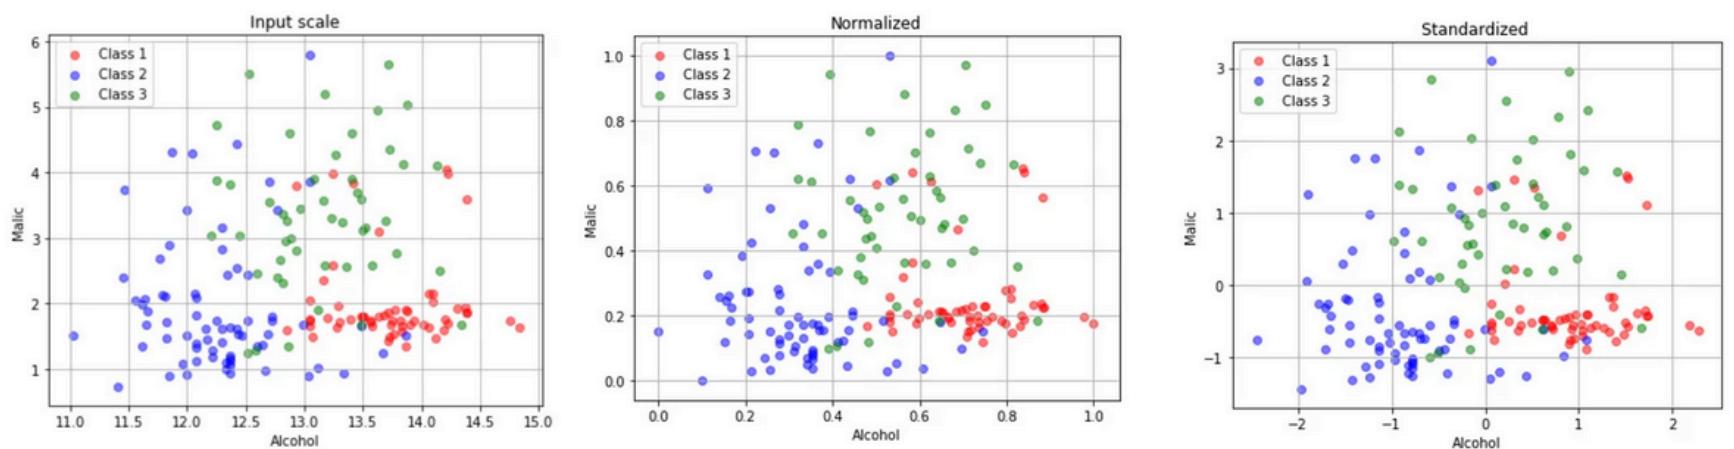
Standardization (Z-score):

- Zero mean, unit variance
- Formula: $z = (x - \mu) / \sigma$

$$z = \frac{x - \mu}{\sigma}$$



Feature Scaling



Feature Scaling

When to use what:

- It depends ;-)
- Standardization:
 - Clustering, PCA
 - Distance-based algorithms
- Normalization:
 - Neural Networks (require data on a 0-1 scale)
 - Image Processing (MNIST Dataset - scale pixel values between 0 and 1 instead of 0 and 255)



31

Let's get our hands dirty

